

# 第十八讲：文件系统实例

## 第 3 节：Zettabyte File System (ZFS)

向勇、陈渝

清华大学计算机系

*xyong,yuchen@tsinghua.edu.cn*

2020 年 4 月 19 日

## 1 第 3 节: Zettabyte File System (ZFS)

- ZFS overview
- ZFS I/O Stack
- ZFS Data Integrity Model

Ref:

- Richard McDougall, Jim Mauro, Solaris Internals: Solaris 10 and OpenSolaris Kernel Architecture, 2nd Edition, Prentice Hall, July 10, 2006, ISBN 0-13-148209-2
- ZFS: The Last Word in File Systems

# What is ZFS?

ZFS is a new kind of filesystem that provides simple administration, transactional semantics, end-to-end data integrity, and immense scalability .

# What is ZFS?

ZFS is a new kind of filesystem that provides simple administration, transactional semantics, end-to-end data integrity, and immense scalability .

- Pooled storage
  - Completely eliminates the antique notion of volumes
  - Does for storage what VM did for memory

# What is ZFS?

ZFS is a new kind of filesystem that provides simple administration, transactional semantics, end-to-end data integrity, and immense scalability .

- Pooled storage
  - Completely eliminates the antique notion of volumes
  - Does for storage what VM did for memory
- Transactional object system
  - Always consistent on disk –no fsck, ever
  - Universal –file, block, iSCSI, swap ...

# What is ZFS?

ZFS is a new kind of filesystem that provides simple administration, transactional semantics, end-to-end data integrity, and immense scalability .

- Pooled storage
  - Completely eliminates the antique notion of volumes
  - Does for storage what VM did for memory
- Transactional object system
  - Always consistent on disk –no fsck, ever
  - Universal –file, block, iSCSI, swap ...
- Provable end-to-end data integrity
  - Detects and corrects silent data corruption
  - Historically considered “too expensive” –no longer true

# What is ZFS?

ZFS is a new kind of filesystem that provides simple administration, transactional semantics, end-to-end data integrity, and immense scalability .

- Pooled storage
  - Completely eliminates the antique notion of volumes
  - Does for storage what VM did for memory
- Transactional object system
  - Always consistent on disk –no fsck, ever
  - Universal –file, block, iSCSI, swap ...
- Provable end-to-end data integrity
  - Detects and corrects silent data corruption
  - Historically considered “too expensive” –no longer true
- Simple administration
  - Concisely express your intent

# ZFS Features

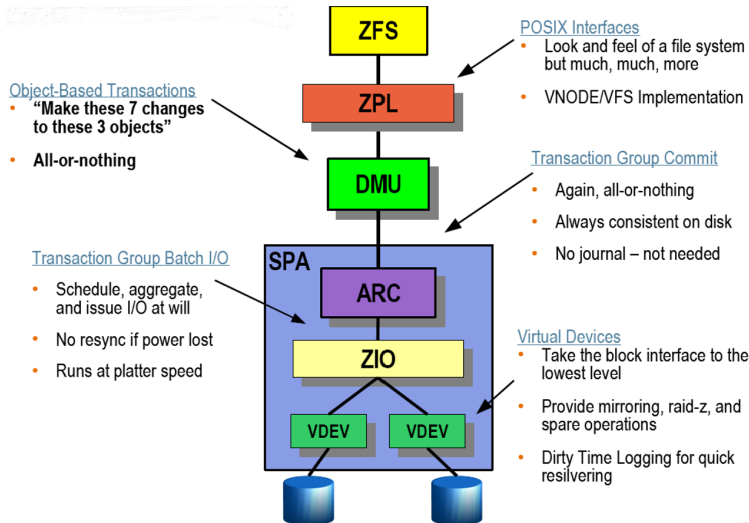
- Immense capacity
  - 128bit
- Provable data integrity
  - Detects and corrects silent data corruption
- Simple administration
  - a pleasure to use



# Pooled storage

- No volume
- Pooled storage
- Many file systems share pool
- And share all I/O channel in the pool


# ZFS I/O Stack



# FS/Volume Model vs. ZFS


## Traditional Volumes

- Abstraction: virtual disk
- Partition/volume for each FS
- Grow/shrink by hand
- Each FS has limited bandwidth
- Storage is fragmented, stranded



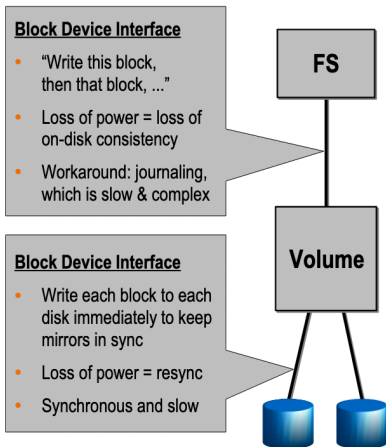
## ZFS Pooled Storage

- Abstraction: malloc/free
- No partitions to manage
- Grow/shrink automatically
- All bandwidth always available
- All storage in the pool is shared

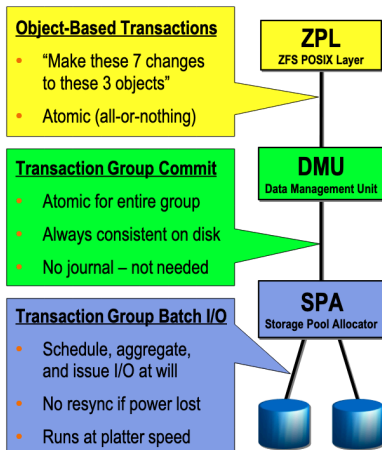


# FS/Volume Interfaces vs. ZFS


## FS/Volume I/O Stack



## ZFS I/O Stack




# ZFL (ZFS POSIX Layer)




- The ZPL is the primary interface for interacting with ZFS as a filesystem.
- It is a layer that sits atop the DMU and presents a filesystem abstraction of files and directories.
- It is responsible for bridging the gap between the VFS interfaces and the underlying DMU interfaces.

# DMU (Data Management Unit)




- Responsible for presenting a transactional object model, built atop the flat address space presented by the SPA.
- Consumers interact with the DMU via objsets, objects, and transactions.
- An objset is a collection of objects, where each object are pieces of storage from the SPA (i.e. a collection of blocks).
- Each transaction is a series of operations that must be committed to disk as a group; it is central to the on-disk consistency for ZFS.

# ARC (Adaptive Replacement Cache)



- DVA (Data Virtual Address) based cache used by DMU
- Self-tuning cache will adjust based on I/O workload
  - > Replaces the page cache
- Central point for memory management for the SPA
  - > Ability to evict buffers as a result of memory pressure


# ZIO (ZFS I/O Pipeline)



- Centralized I/O framework
  - > I/Os follow a structured pipeline
- Translates DVAs to logical locations on vdevs
- Drives dynamic striping and I/O retries across all active vdevs
- Drives compression, checksumming, and data redundancy



# VDEV (Virtual Devices)



- Abstraction of devices
  - > Physical devices (leaf vdevs)
  - > Logical devices (internal vdevs)
- Implementation of data replication algorithms
  - > Mirroring, RAID-Z, and RAID-Z2
- Interfaces with the block level devices
- Provides I/O scheduling and caching

# ZFS Data Integrity Model

- Everything is copy-on-write
  - Never overwrite live data
  - On-disk state always valid –no “windows of vulnerability”
  - No need for fsck(1M)

# ZFS Data Integrity Model


- Everything is copy-on-write
  - Never overwrite live data
  - On-disk state always valid –no “windows of vulnerability”
  - No need for fsck(1M)
- Everything is transactional
  - Related changes succeed or fail as a whole
  - No need for journaling

# ZFS Data Integrity Model


- Everything is copy-on-write
  - Never overwrite live data
  - On-disk state always valid –no “windows of vulnerability”
  - No need for fsck(1M)
- Everything is transactional
  - Related changes succeed or fail as a whole
  - No need for journaling
- Everything is checksummed
  - No silent data corruption

# Copy-On-Write Transactions


1. Initial block tree




2. COW some blocks




3. COW indirect blocks



4. Rewrite uberblock (atomic)



# Constant-Time Snapshots



- At end of TX group, don't free COWed blocks
  - Actually cheaper to take a snapshot than not!
- The tricky part: how do you know when a block is free?

# End-to-End Checksums

## Disk Block Checksums

- Checksum stored with data block
- Any self-consistent block will pass
- Can't even detect stray writes
- Inherent FS/volume interface limitation




Disk checksum only validates media

✓	Bit rot
✗	Phantom writes
✗	Misdirected reads and writes
✗	DMA parity errors
✗	Driver bugs
✗	Accidental overwrite

## ZFS Checksum Trees

- Checksum stored in parent block pointer
- Fault isolation between data and checksum
- Entire pool (block tree) is self-validating




ZFS validates the entire I/O path


✓	Bit rot
✓	Phantom writes
✓	Misdirected reads and writes
✓	DMA parity errors
✓	Driver bugs
✓	Accidental overwrite

# Self-Healing Data in ZFS


1. Application issues a read. ZFS mirror tries the first disk. Checksum reveals that the block is corrupt on disk.



2. ZFS tries the second disk. Checksum indicates that the block is good.



3. ZFS returns good data to the application and repairs the damaged block.





# RAID-Z

Disk		A	B	C	D	E
LBA						
0	P <sub>0</sub>	D <sub>0</sub>	D <sub>2</sub>	D <sub>4</sub>	D <sub>6</sub>	
1	P <sub>1</sub>	D <sub>1</sub>	D <sub>3</sub>	D <sub>5</sub>	D <sub>7</sub>	
2	P <sub>0</sub>	D <sub>0</sub>	D <sub>1</sub>	D <sub>2</sub>		P <sub>0</sub>
3	D <sub>0</sub>	D <sub>1</sub>	D <sub>2</sub>	P <sub>0</sub>	D <sub>0</sub>	
4	P <sub>0</sub>	D <sub>0</sub>	D <sub>4</sub>	D <sub>8</sub>	D <sub>11</sub>	
5	P <sub>1</sub>	D <sub>1</sub>	D <sub>5</sub>	D <sub>9</sub>	D <sub>12</sub>	
6	P <sub>2</sub>	D <sub>2</sub>	D <sub>6</sub>	D <sub>10</sub>	D <sub>13</sub>	
7	P <sub>3</sub>	D <sub>3</sub>	D <sub>7</sub>	P <sub>0</sub>	D <sub>0</sub>	
8	D <sub>1</sub>	D <sub>2</sub>	D <sub>3</sub>	X		P <sub>0</sub>
9	D <sub>0</sub>	D <sub>1</sub>	X	P <sub>0</sub>	D <sub>0</sub>	
10	D <sub>3</sub>	D <sub>6</sub>	D <sub>9</sub>	P <sub>1</sub>	D <sub>1</sub>	
11	D <sub>4</sub>	D <sub>7</sub>	D <sub>10</sub>	P <sub>2</sub>	D <sub>2</sub>	
12	D <sub>5</sub>	D <sub>8</sub>		•	•	•

- Dynamic stripe width
  - Variable block size: 512–128K
  - Each logical block is its own stripe

# RAID-Z

	Disk				
LBA	A	B	C	D	E
0	P <sub>0</sub>	D <sub>0</sub>	D <sub>2</sub>	D <sub>4</sub>	D <sub>6</sub>
1	P <sub>1</sub>	D <sub>1</sub>	D <sub>3</sub>	D <sub>5</sub>	D <sub>7</sub>
2	P <sub>0</sub>	D <sub>0</sub>	D <sub>1</sub>	D <sub>2</sub>	P <sub>0</sub>
3	D <sub>0</sub>	D <sub>1</sub>	D <sub>2</sub>	P <sub>0</sub>	D <sub>0</sub>
4	P <sub>0</sub>	D <sub>0</sub>	D <sub>4</sub>	D <sub>8</sub>	D <sub>11</sub>
5	P <sub>1</sub>	D <sub>1</sub>	D <sub>5</sub>	D <sub>9</sub>	D <sub>12</sub>
6	P <sub>2</sub>	D <sub>2</sub>	D <sub>6</sub>	D <sub>10</sub>	D <sub>13</sub>
7	P <sub>3</sub>	D <sub>3</sub>	D <sub>7</sub>	P <sub>0</sub>	D <sub>0</sub>
8	D <sub>1</sub>	D <sub>2</sub>	D <sub>3</sub>	X	P <sub>0</sub>
9	D <sub>0</sub>	D <sub>1</sub>	X	P <sub>0</sub>	D <sub>0</sub>
10	D <sub>3</sub>	D <sub>6</sub>	D <sub>9</sub>	P <sub>1</sub>	D <sub>1</sub>
11	D <sub>4</sub>	D <sub>7</sub>	D <sub>10</sub>	P <sub>2</sub>	D <sub>2</sub>
12	D <sub>5</sub>	D <sub>8</sub>	•	•	•

- Dynamic stripe width
  - Variable block size: 512–128K
  - Each logical block is its own stripe
- All writes are full-stripe writes
  - Eliminates read-modify-write (it's fast)
  - Eliminates the RAID-5 write hole(no need for NVRAM)

# RAID-Z

	Disk				
LBA	A	B	C	D	E
0	P <sub>0</sub>	D <sub>0</sub>	D <sub>2</sub>	D <sub>4</sub>	D <sub>6</sub>
1	P <sub>1</sub>	D <sub>1</sub>	D <sub>3</sub>	D <sub>5</sub>	D <sub>7</sub>
2	P <sub>0</sub>	D <sub>0</sub>	D <sub>1</sub>	D <sub>2</sub>	P <sub>0</sub>
3	D <sub>0</sub>	D <sub>1</sub>	D <sub>2</sub>	P <sub>0</sub>	D <sub>0</sub>
4	P <sub>0</sub>	D <sub>0</sub>	D <sub>4</sub>	D <sub>8</sub>	D <sub>11</sub>
5	P <sub>1</sub>	D <sub>1</sub>	D <sub>5</sub>	D <sub>9</sub>	D <sub>12</sub>
6	P <sub>2</sub>	D <sub>2</sub>	D <sub>6</sub>	D <sub>10</sub>	D <sub>13</sub>
7	P <sub>3</sub>	D <sub>3</sub>	D <sub>7</sub>	P <sub>0</sub>	D <sub>0</sub>
8	D <sub>1</sub>	D <sub>2</sub>	D <sub>3</sub>	X	P <sub>0</sub>
9	D <sub>0</sub>	D <sub>1</sub>	X	P <sub>0</sub>	D <sub>0</sub>
10	D <sub>3</sub>	D <sub>6</sub>	D <sub>9</sub>	P <sub>1</sub>	D <sub>1</sub>
11	D <sub>4</sub>	D <sub>7</sub>	D <sub>10</sub>	P <sub>2</sub>	D <sub>2</sub>
12	D <sub>5</sub>	D <sub>8</sub>	.	.	.

- Dynamic stripe width
  - Variable block size: 512–128K
  - Each logical block is its own stripe
- All writes are full-stripe writes
  - Eliminates read-modify-write (it's fast)
  - Eliminates the RAID-5 write hole(no need for NVRAM)
- Both single and double parity
- Detects and corrects silent data corruption
  - Checksum-driven combinatorial reconstruction


# RAID-Z

	Disk				
LBA	A	B	C	D	E
0	P <sub>0</sub>	D <sub>0</sub>	D <sub>2</sub>	D <sub>4</sub>	D <sub>6</sub>
1	P <sub>1</sub>	D <sub>1</sub>	D <sub>3</sub>	D <sub>5</sub>	D <sub>7</sub>
2	P <sub>0</sub>	D <sub>0</sub>	D <sub>1</sub>	D <sub>2</sub>	P <sub>0</sub>
3	D <sub>0</sub>	D <sub>1</sub>	D <sub>2</sub>	P <sub>0</sub>	D <sub>0</sub>
4	P <sub>0</sub>	D <sub>0</sub>	D <sub>4</sub>	D <sub>8</sub>	D <sub>11</sub>
5	P <sub>1</sub>	D <sub>1</sub>	D <sub>5</sub>	D <sub>9</sub>	D <sub>12</sub>
6	P <sub>2</sub>	D <sub>2</sub>	D <sub>6</sub>	D <sub>10</sub>	D <sub>13</sub>
7	P <sub>3</sub>	D <sub>3</sub>	D <sub>7</sub>	P <sub>0</sub>	D <sub>0</sub>
8	D <sub>1</sub>	D <sub>2</sub>	D <sub>3</sub>	X	P <sub>0</sub>
9	D <sub>0</sub>	D <sub>1</sub>	X	P <sub>0</sub>	D <sub>0</sub>
10	D <sub>3</sub>	D <sub>6</sub>	D <sub>9</sub>	P <sub>1</sub>	D <sub>1</sub>
11	D <sub>4</sub>	D <sub>7</sub>	D <sub>10</sub>	P <sub>2</sub>	D <sub>2</sub>
12	D <sub>5</sub>	D <sub>8</sub>	.	.	.


- Dynamic stripe width
  - Variable block size: 512–128K
  - Each logical block is its own stripe
- All writes are full-stripe writes
  - Eliminates read-modify-write (it's fast)
  - Eliminates the RAID-5 write hole(no need for NVRAM)
- Both single and double parity
- Detects and corrects silent data corruption
  - Checksum-driven combinatorial reconstruction
- **No special hardware**—ZFS loves cheap disks

# Dynamic Striping

- Writes: striped across all four mirrors
- Reads: wherever the data was written
- Block allocation policy considers:
  - Capacity
  - Performance (latency, BW)
  - Health (degraded mirrors)

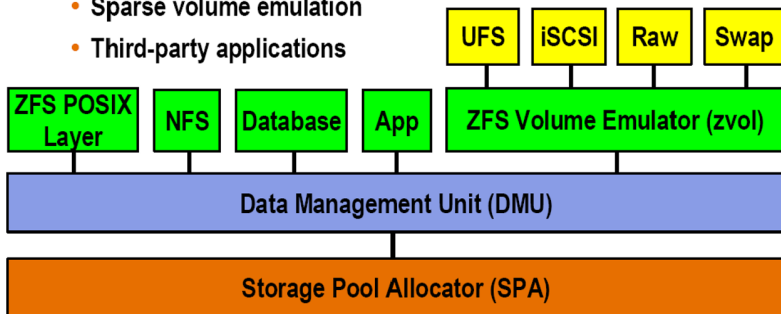


- Writes: striped across all five mirrors
- Reads: wherever the data was written
- No need to migrate existing data
  - Old data striped across 1-4
  - New data striped across 1-5
  - COW gently reallocates old data



# Object-Based Storage

- DMU is a general-purpose transactional object store
  - Filesystems
  - Databases
  - Swap space
  - Sparse volume emulation
  - Third-party applications



# Variable Block Size


- No single block size is optimal for everything
  - Large blocks: less metadata, higher bandwidth
  - Small blocks: more space-efficient for small objects
  - Record-structured files (e.g. databases) have natural granularity; filesystem must match it to avoid read/modify/write

# Variable Block Size

- No single block size is optimal for everything
  - Large blocks: less metadata, higher bandwidth
  - Small blocks: more space-efficient for small objects
  - Record-structured files (e.g. databases) have natural granularity; filesystem must match it to avoid read/modify/write
- Per-object granularity
  - A 37k file consumes 37k –no wasted space
- Enables transparent block-based compression




# Built-in Compression



- Block-level compression in SPA
- Transparent to all other layers

# Built-in Compression



- Block-level compression in SPA
- Transparent to all other layers
- Each block compressed independently
- All-zero blocks converted into file holes
- LZJB and GZIP available today; more on the way